

The full capacity sweep shows how assumed prediction trades memory for tail latency

Projected on the same measured PCIe hierarchy. The 1.5 GiB predictive point reaches 48.0 ms versus 50.6 ms for reactive offload at 6 GiB; all three capacities remain visible.

