

Useful cold experts score higher, but the distributions still overlap

Every distribution scores all nonresident expert IDs in each held-out cold wave. Per-wave standardization removes head-specific logit scale; JS quantifies the class-conditional score divergence (1 bit maximum).

