

The predictors make worse retention choices than ordinary LRU

Observed trace replay at decode, 16 experts resident, 3 layers ahead, and the same one-move-per-wave budget. The right panel explains the left: predictor-guided insertions are less likely than LRU insertions to earn a later hit.

