

Post-training usually changes fewer than one of eight experts at every layer

Observed on exactly matched Base and Instruct tokens. Domain lines expose where the average hides localized substitutions; mathematics at the final layer is the clearest exception.

